


AI Thought Leadership Engine

Hermes CLI · Kanban Orchestration · Multi-Model Cost Optimization

5 Agents · 4 Different Models · Auto-Publish to LinkedIn · ~\$0.006 per post



Why LinkedIn Consistency Is Hard for Educators

Time Drain

Writing one good LinkedIn post takes 1–2 hours.
Multiply by 52 weeks = 104 hours/year lost.

Idea Fatigue

Coming up with fresh, relevant angles every week
becomes mentally exhausting.

Inconsistency


Missing even 2–3 weeks breaks momentum.
Algorithm penalises irregular posters.

The Solution: A fully automated multi-agent pipeline that researches, writes, reviews, and publishes — while you focus on teaching.


WSQ Educator Auto-Publisher — Mohan's LinkedIn Engine

Who	What	Topic	Audience	
Mohan — WSQ trainer, AI courseware author, Singapore tech educator	One high-quality LinkedIn post every Tuesday 08:00 SGT — fully automated	Real trending AI / automation / workforce news tied to Mohan's teaching domain	Singapore HR managers, SME founders, L&D leads, tech-adjacent professionals	
\$0.006 per post	\$0.30 per year	2 hrs saved weekly	52 posts/year	0 min manual effort

End-to-End Pipeline — 5 Agents, 4 Models



Kanban Board — Hermes CLI Stage Tracking



Card moves right as each agent completes →

↑ Card moves to BLOCKED if QA rejects — requires human review before resuming

```
Each card carries:  
week_id · domain · persona · recent_topics[] · post_draft · status · tokens_used · cost_usd
```

Research Agent — Find Real Trending Topics

Role

Find 3 genuinely trending topics with real signals
— never hallucinate events.

Model

Kimi K2 (Moonshot AI)

Why Kimi K2?

- Long context (128K)
- Web grounding built-in
- Cheapest per-token for research tasks
- ~\$0.001 per run

Output

JSON array of 3 topics with signal, audience, angle, confidence

PROMPT

SYSTEM:
You are a content scout for an AI educator in Singapore.
Your only job: find 3 genuinely trending topics from the past 7 days relevant to [DOMAIN]. A topic needs a REAL signal – product launch, viral post, government announcement, research paper, or news event. Never invent signals. If confidence is low, say so.

USER:
Domain: AI agents, no-code automation, workforce transformation, Singapore tech policy

For each topic return JSON with keys:

title	– max 8 words
signal	– what happened and source type
audience	– who on LinkedIn cares
angle	– contrarian how-to case study prediction
confidence	– high medium low

Return a JSON array of exactly 3 topics. Nothing else.

Strategy Agent — Pick the Best Topic Using Reasoning

Role

Receives 3 topics from Scout.
Reasons step-by-step to pick
the best one for this week.

Model

DeepSeek R1 (reasoning)

Why DeepSeek R1?

- Reasoning model — shows chain-of-thought steps
- Better topic ranking than Haiku at 30% lower cost
- ~\$0.002 per run

Output

JSON: chosen topic, reasoning,
hook concept, format, length,
tone

PROMPT

```
SYSTEM:
You are a content strategist. Reason step by step before deciding.
Never guess — show your reasoning chain before the final answer.

USER:
Author: Mohan — WSQ trainer, AI courseware developer, educator
Audience: Singapore HR managers, SME founders, L&D professionals
Voice: direct, practical, educator — not hype, not corporate
Avoid repeating: [RECENT_TOPICS_LIST]
Incoming topics: [SCOUT_JSON]

Reason through EACH question before deciding:
1. Which topic would Mohan have genuine credibility to speak about?
2. Which has strongest engagement for his specific audience?
3. Which has NOT been flooded on LinkedIn this week?
4. Which connects to workforce upskilling or AI adoption in SG?

Output JSON: chosen_topic · why_this_wins · hook_concept
              post_format · target_length · post_tone
```

Writer Agent — Draft the LinkedIn Post (Premium Model Slot)

Role

Write the complete LinkedIn post. This is the highest-quality step in the pipeline.

Model

DeepSeek V3 (best value)

Why DeepSeek V3?

- Writing quality ≈ Sonnet
- 90% cheaper than Sonnet
- Strong English fluency
- Follows style rules well
- ~\$0.002 per run

Hard Rules Enforced

- Never open with 'I'
- No banned words list
- Max 3 hashtags at end
- End with a question

PROMPT

```
SYSTEM:
You are a ghostwriter for a senior AI educator in Singapore.
Write LinkedIn posts that sound like a real person with real opinions –
not a content template, not corporate, not AI-generated.
Reader should finish thinking: 'This person knows their stuff.'
```



```
HARD RULES – violating any is a failure:
- Never open with 'I'
- Never use: "game-changer", "leverage", "delve", "thrilled",
  "excited to share", "the future of", "unlock"
- Maximum 3 hashtags, only at the very end
- End with ONE genuine question that invites a reply
- No em-dash spam (max 1 per post)
- Output ONLY the post text – no preamble, no explanation
```



```
USER:
Topic: [CHOSEN_TOPIC] | Hook: [HOOK_CONCEPT]
Format: [POST_FORMAT] | Tone: [POST_TONE]
Length: [TARGET_LENGTH]
Context: Mohan teaches AI agents to Singapore professionals (WSQ).
```


Quality Assurance Agent — Enforce Rules, Not Style

Role

Mechanical checklist only.
Does NOT rewrite for style.
Only fixes rule violations.

Model

DeepSeek V3 (same as Writer)

Why same model?

- Reuses warm cache
- No cold-start overhead
- Different system prompt = different behaviour
- ~\$0.001 per run

Outputs

approved — no changes
revised — fixed + summary
rejected → card to BLOCKED

PROMPT

```
SYSTEM:
You are a quality checker, NOT a writer.
Enforce rules mechanically. Do not improve style.
Only fix violations — nothing else.

USER:
Run this checklist on the post:
[ ] First word is NOT 'I'
[ ] No banned words: game-changer, leverage, delve,
    "In today's world", thrilled, excited to announce,
    unlock, the future of
[ ] Hashtag count is 3 or fewer, at the end only
[ ] Post ends with a question mark
[ ] Word count within 15% of target: [TARGET_LENGTH]
[ ] No markdown symbols (**, *, #, _)
[ ] Character count under 3000

POST: [DRAFT_POST]

Output JSON: { "status": "approved|revised|rejected",
  "post": "...", "violations_found": [...], "change_summary": "..." }
```

Publisher Agent — Format, Schedule & Trigger LinkedIn API

Role

Mechanical formatter only.
Prepares LinkedIn API payload.
Hermes handles the POST call.

Model

Llama 3.3 70B via Groq

Why Groq + Llama?

- Free tier available
- Ultra-fast inference
- Task is pure formatting
- Zero creativity needed
- Cost: \$0.000

LinkedIn API

POST /v2/ugcPosts
Auth: OAuth 2.0 Bearer
Scheduled: Tue 08:00 SGT

PROMPT

```
SYSTEM:
You are a publishing formatter. Mechanical formatting only.
Do NOT edit content. Do NOT make creative decisions.

USER:
1. If QA status is 'rejected' →
   output { "halt": true, "reason": "QA rejected" } and STOP.

2. Otherwise:
   - Strip all markdown formatting (**, *, #, _)
   - Replace paragraph breaks with \n\n for LinkedIn API
   - Count characters (must be under 3000)
   - Schedule for next Tuesday 08:00 SGT

Output ONLY this JSON (Hermes makes the API call):
{
  "halt": false,
  "text": "...",
  "char_count": 000,
  "word_count": 000,
  "scheduled_for": "YYYY-MM-DDT08:00:00+08:00",
  "platform": "linkedin",
  "ready": true
}
```

Per-Post Cost Breakdown — Right Model for Right Task

Agent	Model	Tokens	Per 1M (in/out)	Per Run
Agent 1 Scout	Kimi K2	~800	\$0.15/\$0.60	~\$0.001
Agent 2 Strategist	DeepSeek R1	~600	\$0.55/\$2.19	~\$0.002
Agent 3 Writer	DeepSeek V3	~1,200	\$0.27/\$1.10	~\$0.002
Agent 4 QA Editor	DeepSeek V3	~800	\$0.27/\$1.10	~\$0.001
Agent 5 Publisher	Groq Llama 3.3	~300	Free tier	\$0.000
TOTAL per post		~3,700 tokens		~\$0.006

Per month (4 posts)

~\$0.025

Per year (52 posts)

~\$0.30

vs All-Claude Sonnet

~\$4.00/yr

Saving vs Sonnet-only

93%

All Models Considered — Why This Stack Wins

Model	Provider	In/1M	Out/1M	Value	Reason	Speed	Best For
Kimi K2	Moonshot	\$0.15	\$0.60	★★★★★	★★★★☆	★★★★★	Research, long-ctx
DeepSeek R1	DeepSeek	\$0.55	\$2.19	★★★★☆	★★★★★	★★★★★	Reasoning, strategy
DeepSeek V3	DeepSeek	\$0.27	\$1.10	★★★★★	★★★★☆	★★★★★	Writing, formatting
Groq Llama 3.3	Groq	Free	Free	★★★★★	★★★☆☆	★★★★★	Simple tasks, speed
Claude Haiku 4.5	Anthropic	\$0.80	\$4.00	★★★☆☆	★★★☆☆	★★★★★	Fallback option
Claude Sonnet 5	Anthropic	\$3.00	\$15.00	★★☆☆☆	★★★★★	★★★☆☆	Not cost-optimal

✓ = Used in pipeline Grey rows = considered but not selected

Fallback Strategy — Pipeline Never Halts

Kimi K2 down → Haiku 4.5 Drop web grounding, use training data only	DeepSeek R1 down → Haiku 4.5 Add chain-of-thought instruction to prompt	DeepSeek V3 down → Haiku 4.5 Slightly higher cost ~\$0.008 total	Groq down → Haiku 4.5 Formatting task only ~\$0.001 extra
---	---	--	---

Hermes auto-routes to fallback — no human intervention needed

Worst case (all providers down): All 5 agents fall back to Claude Haiku 4.5 → Total cost ~\$0.04/post · Still 50× cheaper than manual

Demo Script — 90 Seconds on Stage

00:00	Show Kanban board Open Hermes dashboard. Show INBOX with one card labelled 'Week 32 — AI Agents topic'	00:15	Trigger the pipeline Run: hermes run --card week-32 Watch card move: INBOX → SCOUTING
00:35	Scout + Strategist complete Card moves SCOUTING → DRAFTING. Show the 3 topics returned + which was chosen and why	00:55	Writer completes Card moves DRAFTING → QA. Show the draft post — read the first 2 lines aloud
01:10	QA approves Card moves QA → SCHEDULED. Show status: 'approved', violations_found: []	01:25	Show scheduled payload Show the final JSON payload + cost log: \$0.006 total Say: 'Posts Tuesday 08:00 automatically. Zero minutes of my time.'

Validate Prompts → Build in Hermes CLI → Go Live

VALIDATE BEFORE HERMES CLI

Agent 1	Does Kimi K2 return real signals? Or does it hallucinate events?
Agent 2	Does DeepSeek R1 pick topics that fit YOUR voice and audience?
Agent 3	Does the post sound like YOU? Run 3 topics, compare outputs
Agent 4	Does QA catch violations without over-editing good content?
Agent 5	Does the JSON payload match LinkedIn API spec exactly?

HERMES CLI — BUILD ORDER

Step 1	Test each prompt standalone via Hermes prompt run
Step 2	Wire Agent 1 → 2 handoff test JSON schema passthrough
Step 3	Add Agent 3 + verify tone matches your voice
Step 4	Add Agent 4 QA gate test with intentional violations
Step 5	Wire Agent 5 + LinkedIn OAuth test with a draft post (not publish)
Step 6	Set cron: every Monday 06:00 Go live — first post Tuesday